

CS202: Lab 4: WeensyOS

 cs.nyu.edu/~mwalfish/classes/20sp/labs/lab4.html

[Home](#) | [Schedule](#) | [Policies and grading](#) | [Labs](#) | [Infrastructure](#) | [Exams](#) | [Reference materials](#) | [Announcements](#)


Introduction

In this lab, you will implement process memory isolation, virtual memory, and a system call (`fork()`) in a tiny (but real!) operating system, called WeensyOS.

This will introduce you to virtual memory and reinforce some of the concepts that we have covered this semester.

The WeensyOS kernel runs on x86-64 CPUs. Because the OS kernel runs on the “bare” hardware, debugging kernel code can be tough: if a bug causes misconfiguration of the hardware, the usual reboots won’t help (the applications running on top of it). And because the kernel is the only thing that can access the hardware services (for example, causing the hardware to reboot), it’s a bit tricky to debug what led to a kernel crash can be a bit tricky. One way to develop code for an OS (whether it’s a kernel or a user-space application) was to boot it on a physical CPU. That’s not always the best way to develop code, but it’s a better way since. You will run WeensyOS in QEMU.

QEMU is a software-based virtual machine monitor (VMM) or hypervisor that emulates x86-64 CPU, but if you run it on a real hardware, QEMU itself and the virtual OS you booted and the applications running on it are all in quotation marks because your Linux system is still running on real hardware). So, for example, your last few minutes of work will still get logged to disk (by QEMU running on real hardware), so you don't need to re-running the QEMU emulator application.



Heads up. As always, it's important to start *on time*. In this case, *on time* means 2-3 weeks before the assignment is due, as you will almost certainly need all of the allotted time to complete the lab. Kernel development is less forgiving than developing user-level applications; tiny deviations in the configuration of hardware (such as the MMU) by the OS tend to bring the whole (emulated) machine to a halt.

To save yourself headaches later, read this lab writeup in its entirety before you begin.

Resources.

- You may want to look at Chapter 9 of CS:APP3e (from which our [x86-64 virtual memory handout](#) is borrowed). The book is on reserve at the [Courant library](#). Section 9.7 in particular describes the 64-bit virtual memory architecture of the x86-64 CPU. Figure 9.23 and Section 9.7.1 show and discuss the `PTE_P`, `PTE_W`, and `PTE_U` bits; these are flags in the x86-64 hardware's page table entries that play a central role in this lab.
- You may find yourself during the lab wanting to understand particular assembly instructions. Here are two guides to x86-64 instructions, from [Brown](#) and [CMU](#). The former is more digestible; the latter is more comprehensive. The supplied code also uses certain assembly instructions like `iret`; see [here](#) for a reference.

Getting Started

Obtain the lab files as follows. We assume that you have run the commands in the “Getting Started” section of [lab3](#). To check issue the following command:

```
$ git remote -v
origin  git@github.com:cs202:weensyos.git (fetch)
origin  git@github.com:cs202:weensyos.git (push)
upstream https://github.com:cs202:weensyos.git (fetch)
upstream https://github.com:cs202:weensyos.git (push)
```

The `upstream` should then follow the instructions

Once `$ git remote`

```
$ cd ~/cs202
$ git fetch upstream
$ git merge upstream
```

This lab's files are located

If you have any “conflicts” `push` to save your work



git (fetch)
git (push)

ends in `labs.git`,

g:

urther. Run `git`

Another heads up. Given the complexity of this lab, and the possibility of breaking the functionality of the kernel if you code in some errors, make sure to commit and push your code often! It's very important that your commits have working versions of the code, so if something goes wrong, you can always go back to a previous commit and get back a working copy! At the very least, for this lab, you should be committing once per step (and probably more often), so you can go back to the last step if necessary.

Goal

You will implement complete and correct memory isolation for WeensyOS processes. Then you'll implement full virtual memory, which will improve utilization. You'll implement `fork()` (creating new processes at runtime) and for extra credit, you'll implement `exit()`

(destroying processes at runtime).

We've provided you with a lot of support code for this assignment; the code you will need to write is in fact limited in extent. Our complete solution (for all 5 stages) consists of well under 300 lines of code beyond what we initially hand out to you. All the code you write will go in `kernel.c` (except for part of step 6).

Testing, checking, and validation


For this assignment, your primary checking method will be to run your instance of Weensy OS and visually compare it to the images you see below in the assignment.

Studying these *graphical memory maps* carefully is the best way to determine whether your WeensyOS code for each stage is working correctly. Therefore, you will definitely want to **make sure you understand how to read these maps before you start to code**.

We supply some graphical memory maps for each stage. These will not be your principal source of feedback; they will only tell you whether a given step is passing or failing; look at the test results for more information.

Initial state

Run `make run` in your WeensyOS directory. This will produce a graphical image below, which shows four processes running in parallel. The image is a program in `p-allocator`:



This image loops forever; in an actual run, the bars will move to the right and stay there. Don't worry if your image has different numbers of K's or otherwise has different details.


If your bars run painfully slowly, edit the `p-allocator.c` file and reduce the `ALLOC_SLOWDOWN` constant.

Stop now to read and understand `p-allocator.c`.

Here's how to interpret the memory map display:

- WeensyOS displays the current state of physical and virtual memory. Each character represents 4 KB of memory: a single page. There are 2 MB of physical memory in total. (Ask yourself: how many pages is this?)
- WeensyOS runs four processes, 1 through 4. Each process is compiled from the same source code (`p-allocator.c`), but linked to use a different region of memory.
- Each process asks the kernel for more heap memory, one page at a time, until it runs out of room. As usual, each process's heap begins just above its code and global data, and ends just below its stack. The processes allocate heap memory at different rates: compared to Process 1, Process 2 allocates twice as quickly, Process 3 goes even faster. (A random number generator is used to generate random numbers showing how much heap memory is allocated.)

Here are two labeled memory maps showing how memory is arranged.



The virtual memory display is similar.

- The virtual memory display cycles successively among the four processes' address spaces. In the base version of the WeensyOS code we give you to start from, all four processes' address spaces are the same (your job will be to change that!).
- Blank spaces in the virtual memory display correspond to unmapped addresses. If a process (or the kernel) tries to access such an address, the processor will page fault.
- The character shown at address X in the virtual memory display identifies the owner of the corresponding physical page.
- In the virtual memory display, a character is **reverse video** if an application process is allowed to access the corresponding address. Initially, *any* process can modify *all* of physical memory, including the kernel. Memory is not properly isolated.

Running WeensyOS

Read the `README.md` file for instructions on how to build and run WeensyOS in QEMU's default display causes access to memory outside the console. To make `run-console` the default, you can add the following to your `Makefile`:

There are several ways to add your own code to the WeensyOS kernel. You can add `log_printf` statements to your code, or you can add your own code to the `log.txt` file. You can also use assertions (of the form `assert(condition)`) to check for errors. The helper function `check_page_table_overflow` is provided to check for page table overflow.

Memory system layout

The WeensyOS memory layout is defined by the following constants:

Constant

<code>KERNEL_START_ADDR</code>	Start of kernel code.
<code>KERNEL_STACK_TOP</code>	Top of kernel stack. The kernel stack is one page long.
<code>console</code>	Address of CGA console memory.
<code>PROC_START_ADDR</code>	Start of application code. Applications should not be able to access memory below this address, except for the single page at <code>console</code> .
<code>MEMSIZE_PHYSICAL</code>	Size of physical memory in bytes. WeensyOS does not support physical addresses \geq this value. Defined as <code>0x200000</code> (2MB).
<code>MEMSIZE_VIRTUAL</code>	Size of virtual memory. WeensyOS does not support virtual addresses \geq this value. Defined as <code>0x300000</code> (3MB).



Writing expressions for addresses

WeensyOS uses several C macros to handle addresses. They are defined at the top of `x86-64.h`. The most important include:

Macro	Meaning
<code>PAGESIZE</code>	Size of a memory page. Equals 4096 (or, equivalently, <code>1 << 12</code>).
<code>PAGENUMBER(addr)</code>	Page number for the page containing <code>addr</code> . Expands to an expression analogous to <code>addr / PAGESIZE</code> .
<code>PAGEADDRESS(pn)</code>	The initial address (zeroth byte) in page number <code>pn</code> . Expands to an expression analogous to <code>pn * PAGESIZE</code> .
<code>PAGEINDEX(addr, level)</code>	The index in the <code>level</code> th page table for <code>addr</code> . <code>level</code> must be between 0 and 3; 0 returns the level-1 page table index (address bits 39–47), 1 returns the level-2 index (bits 30–38), 2 returns the level-3 index (bits 21–29), and 3 returns
<code>PTE_ADDR(pe)</code>	Physical address of the page table entry <code>pe</code> . The low-order

Before you begin coding, you must know how the macros represent and be able to derive values from them.

Kernel and process

The version of WeensyOS we are using has a single kernel and all processes in a single, shared address space. This is defined by the `kernel_pagetable` page table. The mapping of virtual address `X` maps to physical address `Y` is defined by the *identity mapping*:

As you work through the chapter, you will see how independent address spaces, where each process can access only a subset of physical memory.

The kernel, though, must remain able to access *any* location in physical memory. Therefore, all kernel functions run using the `kernel_pagetable` page table. Thus, in kernel functions, each virtual address maps to the physical address with the same number. The `exception()` function explicitly installs `kernel_pagetable` when it begins.

WeensyOS system calls are more expensive than they need to be, since every system call switches address spaces twice (once to `kernel_pagetable` and once back to the process's page table). Real-world operating systems avoid this overhead. To do so, real-world kernels access memory using *process* page tables, rather than a kernel-specific `kernel_pagetable`. This makes a kernel's code more complicated, since kernels can't always access all of physical memory directly under that design.